

5SENG003W - Algorithms, Week 9

Dr. Klaus Draeger

RECAP: Graphs

- ▶ Graphs consist of **vertices** connected by **edges**
- ▶ Edges can be **directed** or **undirected**
- ▶ Representations:
 - ▶ Set of vertex pairs
 - ▶ Ordered pairs for directed graphs:
 $\{(v_1, v_2), (v_1, v_3), (v_2, v_5), \dots\}$
 - ▶ Unordered pairs for undirected graphs:
 $\{\{v_1, v_7\}, \{v_2, v_5\}, \{v_2, v_3\}, \dots\}$
 - ▶ Adjacency matrix
 - ▶ $a[i][j]$ is 1 if there is an edge (v_i, v_j) and 0 otherwise
 - ▶ Adjacency lists
 - ▶ Each vertex contains a list of outgoing edges
 - ▶ Vertices can be implicit: can represent the graph by an array of lists

Recap: Graph properties

- ▶ **Adjacency** between vertices, **incidence** between vertices and edges
- ▶ (In-, out-) degree
- ▶ Walks, paths, cycles
- ▶ Loops and parallel edges
- ▶ Connectedness, acyclicity
- ▶ Colours, weights, ...

More about graph representations

- ▶ We have previously seen that graphs can be represented using
 - ▶ Adjacency **matrices**
 - ▶ Adjacency **lists**
- ▶ Let us talk about this in more detail. The main questions to ask yourself are
 - ▶ How do we represent **vertices**?
 - ▶ How do we represent **edges**?
 - ▶ how do we store **additional data** (labels, weights, ...)?

Representing vertices

- ▶ We can represent vertices using
 - ▶ Explicit **Vertex** objects
 - ▶ Containing the adjacency list (makes no sense with an adjacency matrix)
 - ▶ Containing all relevant additional data
 - ▶ Suitable if number of vertices is initially unknown
 - ▶ IDs
 - ▶ To index the adjacency matrix row or the adjacency list (in the array of lists)
 - ▶ Could be more than a single number (e.g. co-ordinates on a chessboard)

Representing edges

- ▶ We can represent edges using
 - ▶ Explicit **Edge** objects
 - ▶ Containing the target vertex
 - ▶ Containing all relevant additional data
 - ▶ Suitable if there are parallel edges or complex additional data
 - ▶ Just the target vertex
 - ▶ ID or pointer to object

Recap: Searching in graphs

- ▶ A common operation on graphs is **exploration**
 - ▶ Searching for some particular **vertex**
 - ▶ Searching for a **path** between two vertices
 - ▶ Searching for a **shortest** path between two vertices
 - ▶ Or just trying to **fully explore** the graph
- ▶ Graph exploration is more complex than in trees or lists
 - ▶ Graphs can have **cycles**
 - ▶ Naive search could get stuck in a loop
 - ▶ Need to keep track of visited vertices to prevent this
- ▶ Previously introduced: **depth-first** and **breadth-first** search

Recap: Depth-first search

- ▶ In depth-first search, we recursively follow outgoing edges
- ▶ To explore a vertex, we
 - ▶ Pick a previously unvisited neighbour
 - ▶ Recursively explore that neighbour
 - ▶ Repeat if (after returning from the recursive call) there are still unvisited neighbours
- ▶ Note the difference between **visited** and **explored**
 - ▶ A node is **visited** as soon as we encounter it
 - ▶ A node is **explored** once we have visited all its neighbours (possibly from a different node)
 - ▶ This difference will be represented by the **closed** and **open** lists

Recap: Breadth-first search

- ▶ In breadth-first search, we explore the graph in "layers"
- ▶ Not recursive
- ▶ We maintain two data structures:
 - ▶ The set of all visited vertices (the **closed list**)
 - ▶ A queue of visited but not yet fully explored vertices (the **open list**)
- ▶ In each iteration, we
 - ▶ Check the neighbours of the open list's front element
 - ▶ Enqueue those which were not yet visited (and add them to the closed list)
 - ▶ Dequeue the front element (it is now fully explored)

Recap: Iterative depth-first search

- ▶ We can perform depth-first search in the same iterative style as breadth-first search
- ▶ The key change is organising the open list as a **stack** rather than a queue
- ▶ In each iteration, we
 - ▶ Check if the top element has unvisited neighbours
 - ▶ If yes, push one of them onto the open list (and add it to the closed list)
 - ▶ If not, pop the top element off (it is now fully explored)
- ▶ Alternatively, in each iteration, we
 - ▶ Pop the top element of the stack
 - ▶ Push **all** its unvisited neighbours onto the open list (and add them to the closed list)
- ▶ The first version means re-examining adjacency lists but will directly produce a path to the target (if any)

Depth-first search: iterative version

- ▶ The second version of iterative DFS looks like this (assuming for simplicity that Vertex contains

List<Vertex> adjacencyList;)

```
public void DFS(Vertex source) {  
    // Set up open and closed list  
    Stack<Vertex> o = new Stack<Vertex>();  
    HashSet<Vertex> c = new HashSet<Vertex>();  
    o.push(source);  
    c.add(source);  
    while(!o.empty()) {  
        Vertex v = o.pop();  
        for(Vertex w : v.adjacencyList)  
            if(! c.contains(w)) { // not yet visited  
                o.push(w);  
                c.add(w);  
            }  
    }  
}
```

Aside: eliminating recursion

- ▶ Iterative DFS is an example of a more general idea: We can get rid of recursion by using a stack (to store pending branches)
- ▶ Example: pre-order traversal in a tree looks like this:

```
public class TreeNode{  
    public int data;  
    public TreeNode leftChild, rightChild;  
}  
  
public void preOrder(TreeNode n){  
    if(n != null){  
        System.out.println(n.data);  
        preOrder(n.leftChild);  
        preOrder(n.rightChild);  
    }  
}
```

Aside: eliminating recursion

- ▶ Example: an **iterative** version of pre-order traversal (using java.util.Stack)

```
public class TreeNode{
    public int data;
    public TreeNode leftChild, rightChild;
}

public void preOrder(TreeNode t){
    Stack<TreeNode> s = new Stack<TreeNode>();
    s.push(t);
    while(!s.empty()){
        // pop top branch off the stack
        t = s.pop();
        if(t != null){ // output data and push children on the stack
            System.out.println(t.data);
            s.push(t.rightChild);
            s.push(t.leftChild);
        }
    }
}
```

- ▶ This code could still be improved (some vertices get pushed and then immediately popped)
- ▶ Think about how to do this with in-order or post-order

Exploration for pathfinding

- ▶ The traversal algorithms explore the entire graph
- ▶ When searching for a path to a target this is not needed. Instead we can
 - ▶ Abort the search as soon as the target is found
 - ▶ Keep track of the path taken, e.g. using a **predecessor map**
 - ▶ This map represents a **spanning tree** with predecessors being parents
- ▶ Either algorithm will find a path if one exists
- ▶ BFS also guarantees that the path found has **minimal length**

Why breadth-first search works


- ▶ BFS advances in "layers":
 - ▶ For a vertex v , let $f(v)$ be the distance from the start to v .
 - ▶ BFS first visits first all v with $f(v) = 0$, then $f(v) = 1$ etc
- ▶ Why is this the case?
 - ▶ At first, the open list only contains the start vertex, which is the only vertex s with $f(s) = 0$
 - ▶ All vertices with $f(v) = k + 1$ are reached in one step from vertices with $f(v) = k$
 - ▶ So they are all enqueued while the vertices with $f(v) = k$ are getting visited

Breadth-first search breaks for weighted graphs

- ▶ In a weighted graph, vertices are no longer found ordered by distance
- ▶ For example, vertices with $f(v) = 5$ may have edges to vertices with $f(v) = k$ for **any** $k > 5$.
- ▶ Idea: could use a **sorted** data structure like a binary search tree for the open list to fix this
 - ▶ Then insertion and extraction in the open list both are in $O(\log n)$
- ▶ But we only ever need the minimal-distance vertex
 - ▶ So we can use a **priority queue** based on a **min-heap**.
- ▶ This gives us **Dijkstra's Algorithm**.
 - ▶ Finds shortest paths from the start to all other targets
 - ▶ Easy to turn into a single-target version

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_0
- Closed list:
 $\{v_0 \mapsto 0\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_0 \ v_1$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_0 v_1 v_5
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_5 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_1 \ v_5$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_5 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_1 \ v_2 \ v_5$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_5 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_1 \ v_2 \ v_5$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_2 v_5
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_2 \ v_3 \ v_5$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_3 v_5
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_3 v_5 v_4
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_5 v_4
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 $v_5 \ v_4 \ v_6$
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4, v_6 \mapsto 9\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_4 v_6
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4, v_6 \mapsto 9\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_4 v_6
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4, v_6 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example


- Consider this example:



- Open list:
 v_6
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4, v_6 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm: example

- Consider this example:



- Open list:
- Closed list:
 $\{v_0 \mapsto 0, v_1 \mapsto 1, v_2 \mapsto 2, v_3 \mapsto 3, v_4 \mapsto 4, v_5 \mapsto 4, v_6 \mapsto 5\}$
- Step 5: Vertex v_2 inserted in front of v_5
- Done; found shortest distances to all vertices

Dijkstra's algorithm in a nutshell

- ▶ In a nutshell, Dijkstra's algorithm does the following:
 - ▶ For each vertex v , keep track of the shortest known distance $d(s, v)$ from s to v
 - ▶ Many versions of the algorithm initially set all distances to ∞ except s , where it is 0
 - ▶ Our **closed list** represents those vertices where we have found a non- ∞ distance
 - ▶ In each iteration, expand a vertex v
 - ▶ This means exploring its outgoing edges (v, x)
 - ▶ v must be a previously unexpanded vertex with minimal $d(s, v)$
 - ▶ Our **open list** contains the candidates for this
 - ▶ During the expansion, update $d(s, x)$ if $d(s, v) + w(v, x)$ is smaller

Dijkstra's algorithm: details

- ▶ We want to be able to output actual paths
- ▶ So we keep track of the **predecessor** of each vertex (i.e. from which other vertex we reached it, getting the shortest distance)
- ▶ Final output in the example:

| Vertex | Distance | Shortest path |
|--------|----------|--------------------------------|
| v_0 | 0 | v_0 |
| v_1 | 1 | v_0, v_1 |
| v_2 | 2 | v_0, v_1, v_2 |
| v_3 | 3 | v_0, v_1, v_2, v_3 |
| v_4 | 4 | v_0, v_1, v_2, v_3, v_4 |
| v_5 | 4 | v_0, v_1, v_5 |
| v_6 | 5 | $v_0, v_1, v_2, v_3, v_4, v_6$ |

Dijkstra's algorithm: implementation

- ▶ We make a few changes to our existing code:
- ▶ We use a dedicated **Vertex** class
 - ▶ This allows us to store additional information in them, which will be helpful soon
- ▶ We use maps (**java.util.HashMap** or **std::map**)
 - ▶ In adjacency lists to keep track of weights

```
public class Vertex{  
    public int id;  
  
    public HashMap<Vertex, Integer> adjacencyList;  
  
    /* ... */  
}
```

- ▶ In the closed list to keep track of shortest path lengths

Dijkstra's algorithm: implementation

- ▶ We make a few changes to our existing code:
- ▶ We use a wrapper class containing
 - ▶ A vertex
 - ▶ The length of the shortest known path from s to v

```
public class WrappedVertex{  
    public Vertex vertex;  
  
    // shortest currently known distance from the start  
    public int distance;  
  
    public WrappedVertex(Vertex v, int d){  
        vertex = v;  
        distance = d;  
    }  
}
```

in order to store vertices in a heap

Dijkstra's algorithm: implementation

- ▶ We adapt the MinHeap class
 - ▶ It now stores wrapped vertices
 - ▶ It uses the current distance to figure out how far to sift up a new entry

```
public class MinHeap {
    ArrayList<WrappedVertex> items;

    public MinHeap() {
        items = new ArrayList<WrappedVertex>();
    }

    public void insert(Vertex v, int d) {
        WrappedVertex newItem = new WrappedVertex(v, d);
        int index = items.size();
        items.add(newItem);
        // Determine position for insertion
        while(index > 0) {
            int parent = (index-1)/2; // parent index
            if(items.get(parent).distance > newItem.distance) { // sift up
                items.set(index, items.get(parent));
                index = parent;
            }
            else
                break;
        }
        items.set(index, newItem);
    }

    /* ... */
}
```

Dijkstra's algorithm: implementation

- ▶ Finally, we adapt the search method we have seen before to use a heap as its open list
- ▶ Note that if we find a new, shorter path to a vertex v , we
 - ▶ Create a new WrappedVertex with the new distance
 - ▶ Insert it in the open list
 - ▶ We may have several wrapped versions of the same vertex
 - ▶ They will have different distances
 - ▶ The one with the smallest distance will be handled first
 - ▶ The others will just be discarded eventually

Dijkstra's algorithm: implementation

- The structure of the main loop looks like this:

```
public void Dijkstra(Vertex source){
    // Set up open, closed, and predecessor lists
    MinHeap o = new MinHeap();
    HashMap<Vertex, Integer> c = new HashMap<>();
    HashMap<Vertex, Vertex> p = new HashMap<>();
    o.insert(source, 0);
    c.put(source, 0);
    p.put(source, null);
    while(!o.empty()){
        WrappedVertex w = o.extractMinimum();
        if(w.distance == c.get(w.vertex)){ // w is up to date
            /* TO DO: complete this
               Output the distance and shortest path (follow the chain of predecessors)
               For every edge e out of w.vertex,
                   - check if e gives a new shortest path to its target
                   - update lists if yes
            */
        }
    }
}
```

Dijkstra's algorithm: analysis

- ▶ Dijkstra's algorithm is an example of a **greedy** algorithm
 - ▶ It always picks the option that currently looks best
- ▶ It expands each vertex (i.e. examines its outgoing edges) **once**
- ▶ Its complexity depends on details such as graph representation, sparsity, etc.
A general lower bound is $\Theta(|E| + |V|^2)$.