

5SENG001W - Algorithms, Week 10


Dr. Klaus Draeger

Recap: Dijkstra's algorithm

- ▶ Dijkstra's algorithm does the following:
 - ▶ For each vertex v , keep track of the shortest known distance $d(s, v)$ from s to v
 - ▶ Many versions of the algorithm initially set all distances to ∞ except s , where it is 0
 - ▶ Our **closed list** represents those vertices where we have found a non- ∞ distance
 - ▶ In each iteration, expand a vertex v
 - ▶ This means exploring its outgoing edges (v, x)
 - ▶ v must be a previously unexpanded vertex with minimal $d(s, v)$
 - ▶ Our **open list** contains the candidates for this
 - ▶ During the expansion, update $d(s, x)$ if $d(s, v) + w(v, x)$ is smaller

Towards A*


- Suppose we want to get from S to F in this graph:



- Why do we bother exploring to the south and west?

Towards A*


- Suppose we want to get from S to F in this graph:



- Why do we bother exploring to the south and west?

Towards A*


- Suppose we want to get from S to F in this graph:



- Why do we bother exploring to the south and west?

Towards A^*

- Suppose we want to get from S to F in this graph:




- Why do we bother exploring to the south and west?

Towards A*

- ▶ In the grid example, we know exactly where to go
 - ▶ We have a notion of directions, and know in which direction the target is
 - ▶ We know **how far** any given vertex is from the target
- ▶ So pathfinding becomes trivial
- ▶ In general, we don't have this exact information
 - ▶ We would need to already have solved the path-finding problem
- ▶ But we may have a suitable **estimate** of distances to the target

Towards A*

- For example, if the edges in the grid



have unknown weights on them

- The edges on the direct paths may have higher weights
- So it may make sense to take a detour
- But the distance in the unweighted version (“as the crow flies”) can still be used to steer the search
- We use it to bias the exploration to vertices that are “in the right direction”

The A* algorithm

- ▶ Like Dijkstra's algorithm, A* organises its open list as a priority queue
- ▶ The value associated with a vertex v is $f(v) + h(v)$, where
 - ▶ $f(v)$ is the distance from the start to v
(discovered during the exploration)
 - ▶ $h(v)$ is an estimate for the distance from v to the target
(must be provided)

Estimate functions


- ▶ A* relies on an estimate function h to work.
- ▶ This function should have some properties in order to work well:
 - ▶ It needs to be cheap to compute for each vertex
 - ▶ It needs to be an **underapproximation**, i.e. $h(v)$ cannot be greater than the actual distance to the target (otherwise a suboptimal path might be found first)
- ▶ How to find such an estimate?

State graphs and relaxations

- ▶ Common case: the graph represents some **system**
 - ▶ Vertices are **states** of the system
 - ▶ Edges are **transitions**
 - ▶ We want to go from some initial to some final state
- ▶ This graph is potentially huge
 - ▶ We do **not** want to create the whole graph before looking for a path
 - ▶ Instead, we want to create vertices only when needed ("on the fly")

State graphs: example

- ▶ We previously saw this tile sliding puzzle
 - ▶ Eight tiles numbered 1, ..., 8 in a 3*3 square (bigger versions also exist)
 - ▶ One hole which adjacent tiles can be moved into
 - ▶ Goal is to have all the tiles in order



State graphs: example

- ▶ The state graph for the 9-puzzle has
 - ▶ One vertex for each permutation of the tiles and hole
 - ▶ Edges corresponding to moves in the puzzle
- ▶ A total of $9! = 362880$ vertices
(and $16! = 20,922,789,888,000$ for the 4×4 version)
so we really don't want to construct the whole graph
- ▶ On-the-fly search using A^* helps avoid this

State graphs and relaxations

- ▶ Estimate functions for state graphs can be found based on **relaxations**
- ▶ Idea: relax the interaction between parts, so that
 - ▶ They can move independently
(so it is easier to compute the distance in the relaxation)
 - ▶ All moves in the original system are still in the relaxation
(so the distance in the relaxation is an underapproximation)
- ▶ Let's have a look at how this works in the 9-puzzle

State graphs: example

- ▶ In the 9-puzzle, the parts are the numbered tiles
- ▶ Their moves are not independent
(because they cannot share a position)
- ▶ We can relax the system by removing this constraint
- ▶ In the relaxation, we can move each tile to its place independently
- ▶ The distance in the relaxation is the sum of each tile's distance to its target position

State graphs: example

- ▶ For example, suppose we start in this state s_1 :


1	3	2
8	5	4
	6	7

gets the distance estimate $h(s_1) = 10$, because

- ▶ Tiles 1 and 5 are in their target position
- ▶ Tiles 2,3 are 1 step away from their target location
- ▶ Tiles 4,6,7,8 are 2 steps away from their target location
- ▶ $0 + 1 + 1 + 2 + 0 + 2 + 2 + 2 = 10$

State graphs: example

- ▶ The estimates for the successor states are 9 and 11:



- ▶ Then the search progresses like this:
 - ▶ The distances from the start are $f(s_1) = 0, f(s_2) = f(s_3) = 1$.
 - ▶ The priorities are $f(s_1) + h(s_1) = 10, f(s_2) + h(s_2) = 12, f(s_3) + h(s_3) = 10$.
 - ▶ So s_3 comes before s_2 in the priority queue.